

Implement SGD Classifier with Logloss and L2 regularization Using SGD without using sklearn

There will be some functions that start with the word "grader" ex: grader_weights(), grader_sigmoid(), grader_logloss() etc, you should not change those function definition.

Every Grader function has to return True.

Importing packages

```
In [365]: import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn import linear_model
```

Creating custom dataset

```
In [366]: # please don't change random_state
X, y = make_classification(n_samples=50000, n_features=15, n_informative=10, n_re
                                n_classes=2, weights=[0.7], class_sep=0.7, random_stat
# make_classification is used to create custom dataset
# Please check this Link (https://scikit-learn.org/stable/modules/generated/sklearn
```

```
In [367]: X.shape, y.shape
```

```
Out[367]: ((50000, 15), (50000,))
```

Splitting data into train and test

```
In [368]: #please don't change random_state
# you need not standardize the data as it is already standardized
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_
```

```
In [369]: X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

```
Out[369]: ((37500, 15), (37500,), (12500, 15), (12500,))
```

SGD classifier

```
In [370]: # alpha : float
# Constant that multiplies the regularization term.

# eta0 : double
# The initial learning rate for the 'constant', 'invscaling' or 'adaptive' schedu

clf = linear_model.SGDClassifier(eta0=0.0001, alpha=0.0001, loss='log', random_st
clf
# Please check this documentation (https://scikit-learn.org/stable/modules/genera
```

```
Out[370]: SGDClassifier(eta0=0.0001, learning_rate='constant', loss='log',
                      random_state=15, verbose=2)
```

```
In [371]: clf.fit(X=X_train, y=y_train) # fitting our model
```

```
-- Epoch 1
Norm: 0.77, NNZs: 15, Bias: -0.316653, T: 37500, Avg. loss: 0.455552
Total training time: 0.01 seconds.
-- Epoch 2
Norm: 0.91, NNZs: 15, Bias: -0.472747, T: 75000, Avg. loss: 0.394686
Total training time: 0.02 seconds.
-- Epoch 3
Norm: 0.98, NNZs: 15, Bias: -0.580082, T: 112500, Avg. loss: 0.385711
Total training time: 0.03 seconds.
-- Epoch 4
Norm: 1.02, NNZs: 15, Bias: -0.658292, T: 150000, Avg. loss: 0.382083
Total training time: 0.04 seconds.
-- Epoch 5
Norm: 1.04, NNZs: 15, Bias: -0.719528, T: 187500, Avg. loss: 0.380486
Total training time: 0.05 seconds.
-- Epoch 6
Norm: 1.05, NNZs: 15, Bias: -0.763409, T: 225000, Avg. loss: 0.379578
Total training time: 0.05 seconds.
-- Epoch 7
Norm: 1.06, NNZs: 15, Bias: -0.795106, T: 262500, Avg. loss: 0.379150
Total training time: 0.06 seconds.
-- Epoch 8
Norm: 1.06, NNZs: 15, Bias: -0.819925, T: 300000, Avg. loss: 0.378856
Total training time: 0.07 seconds.
-- Epoch 9
Norm: 1.07, NNZs: 15, Bias: -0.837805, T: 337500, Avg. loss: 0.378585
Total training time: 0.08 seconds.
-- Epoch 10
Norm: 1.08, NNZs: 15, Bias: -0.853138, T: 375000, Avg. loss: 0.378630
Total training time: 0.09 seconds.
Convergence after 10 epochs took 0.09 seconds
```

```
Out[371]: SGDClassifier(eta0=0.0001, learning_rate='constant', loss='log',
                      random_state=15, verbose=2)
```

```
In [372]: clf.coef_, clf.coef_.shape, clf.intercept_
#clf.coef_ will return the weights
#clf.coef_.shape will return the shape of weights
#clf.intercept_ will return the intercept term
```

```
Out[372]: (array([[ -0.42336692,  0.18547565, -0.14859036,  0.34144407, -0.2081867 ,
                    0.56016579, -0.45242483, -0.09408813,  0.2092732 ,  0.18084126,
                    0.19705191,  0.00421916, -0.0796037 ,  0.33852802,  0.02266721]]),
          (1, 15),
          array([-0.8531383]))
```

Implement Logistic Regression with L2 regularization Using SGD: without using sklearn

1. We will be giving you some functions, please write code in that functions only.
2. After every function, we will be giving you expected output, please make sure that you get that output.

- Initialize the weight_vector and intercept term to zeros (Write your code in `def initialize_weights()`)
- Create a loss function (Write your code in `def logloss()`)

$$\text{logloss} = -1 * \frac{1}{n} \sum_{\text{foreach } Y_t, Y_{\text{pred}}} (Y_t \log_{10}(Y_{\text{pred}}) + (1 - Y_t) \log_{10}(1 - Y_{\text{pred}}))$$

- for each epoch:
 - for each batch of data points in train: (keep batch size=1)
 - calculate the gradient of loss function w.r.t each weight in weight vector (write your code in `def gradient_dw()`)

$$dw^{(t)} = x_n(y_n - \sigma((w^{(t)})^T x_n + b^t)) - \frac{\lambda}{N} w^{(t)}$$

- Calculate the gradient of the intercept (write your code in `def gradient_db()`) [check this \(https://drive.google.com/file/d/1nQ08-XY4zvOLzRX-IGf8EYB5arb7-m1H/view?usp=sharing\)](https://drive.google.com/file/d/1nQ08-XY4zvOLzRX-IGf8EYB5arb7-m1H/view?usp=sharing)

$$db^{(t)} = y_n - \sigma((w^{(t)})^T x_n + b^t)$$

- Update weights and intercept (check the equation number 32 in the above mentioned pdf (<https://drive.google.com/file/d/1nQ08-XY4zvOLzRX-IGf8EYB5arb7-m1H/view?usp=sharing>)):

$$w^{(t+1)} \leftarrow w^{(t)} + \alpha(dw^{(t)})$$

$$b^{(t+1)} \leftarrow b^{(t)} + \alpha(db^{(t)})$$

- calculate the log loss for train and test with the updated weights (you can check the python assignment 10th question)
- And if you wish, you can compare the previous loss and the current loss, if it is not updating, then you can stop the training

- append this loss in the list (this will be used to see how loss is changing for each epoch after the training is over)

Initialize weights

```
In [373]: def initialize_weights(row_vector):
    ''' In this function, we will initialize our weights and bias'''
    #initialize the weights as 1d array consisting of all zeros similar to the di
    #zeros_like function to initialize zero to weights
    #https://docs.scipy.org/doc/numpy/reference/generated/numpy.zeros_like.html
    w = np.zeros_like(X_train[0])
    #initialize bias to zero
    b = 0
    return w,b
```

```
In [374]: dim=X_train[0]
row_vector = len(dim)
w,b = initialize_weights(row_vector)
print('w =',(w))
print('b =',str(b))
```

```
w = [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
b = 0
```

Grader function - 1

```
In [375]: dim=X_train[0]
w,b = initialize_weights(dim)
def grader_weights(w,b):
    assert((len(w)==len(dim)) and b==0 and np.sum(w)==0.0)
    return True
grader_weights(w,b)
```

Out[375]: True

Compute sigmoid

$$\text{sigmoid}(z) = 1/(1 + \exp(-z))$$

```
In [376]: def sigmoid(z):
    ''' In this function, we will return sigmoid of z'''
    # compute sigmoid(z) and return
    return 1/(1+np.exp(-z))
```

Grader function - 2

```
In [377]: def grader_sigmoid(z):
            val=sigmoid(z)
            assert(val==0.8807970779778823)
            return True
            grader_sigmoid(2)
```

Out[377]: True

Compute loss

$$\text{logloss} = -1 * \frac{1}{n} \sum_{\text{foreach } Y_t, Y_{\text{pred}}} (Y_t \log_{10}(Y_{\text{pred}}) + (1 - Y_t) \log_{10}(1 - Y_{\text{pred}}))$$

```
In [378]: def compute_log_loss(y_true,y_pred):
            #Number of rows n
            sum = 0 # initializing sum as zero
            for i in range (len(y_true)):
                #using the formula for f(Y,Yscore)
                sum += ((y_true[i] * np.log10(y_pred[i])) + ((1-y_true[i]) * (np.log10(1-
            # Computing log loss
            loss = -(sum/n)
            return loss
```

Grader function - 3

```
In [379]: #round off the value to 8 values
            def grader_logloss(true,pred):
                loss=logloss(true,pred)
                assert(np.round(loss,6)==0.076449)
                return True
            true=np.array([1,1,0,1,0])
            pred=np.array([0.9,0.8,0.1,0.8,0.2])
            grader_logloss(true,pred)
```

Out[379]: True

Compute gradient w.r.to 'w'

$$dw^{(t)} = x_n (y_n - \sigma((w^{(t)})^T x_n + b^t)) - \frac{\lambda}{N} w^{(t)}$$

```
In [380]: #make sure that the sigmoid function returns a scalar value, you can use dot func
            def gradient_dw(x,y,w,b,alpha,N):
                '''In this function, we will compute the gardient w.r.to w '''

                #Using the above equation for gradient

                dw = x * (y-sigmoid(np.dot(w,x)+b)-(alpha/N)*w)

                return dw
```

Grader function - 4

```
In [381]: def grader_dw(x,y,w,b,alpha,N):
            grad_dw=gradient_dw(x,y,w,b,alpha,N)
            assert(np.round(np.sum(grad_dw),5)==4.75684)
            return True
            grad_x=np.array([-2.07864835,  3.31604252, -0.79104357, -3.87045546, -1.14783286,
                             -2.81434437, -0.86771071, -0.04073287,  0.84827878,  1.99451725,
                             3.67152472,  0.01451875,  2.01062888,  0.07373904, -5.54586092])
            grad_y=0
            grad_w=np.array([ 0.03364887,  0.03612727,  0.02786927,  0.08547455, -0.12870234,
                             -0.02555288,  0.11858013,  0.13305576,  0.07310204,  0.15149245,
                             -0.05708987, -0.064768 ,  0.18012332, -0.16880843, -0.27079877])
            grad_b=0.5
            alpha=0.0001
            N=len(X_train)
            grader_dw(grad_x,grad_y,grad_w,grad_b,alpha,N)
```

Out[381]: True

Compute gradient w.r.to 'b'

$$db^{(t)} = y_n - \sigma((w^{(t)})^T x_n + b^t)$$

```
In [382]: #sb should be a scalar value
            def gradient_db(x,y,w,b):
                '''In this function, we will compute gradient w.r.to b '''
                #Using the above equation
                db = y - sigmoid(np.dot(w,x)+b)

                return db
```

Grader function - 5

```
In [383]: def grader_db(x,y,w,b):
            grad_db=gradient_db(x,y,w,b)
            assert(np.round(grad_db,4)==-0.3714)
            return True
            grad_x=np.array([-2.07864835,  3.31604252, -0.79104357, -3.87045546, -1.14783286,
                             -2.81434437, -0.86771071, -0.04073287,  0.84827878,  1.99451725,
                             3.67152472,  0.01451875,  2.01062888,  0.07373904, -5.54586092])
            grad_y=0.5
            grad_b=0.1
            grad_w=np.array([ 0.03364887,  0.03612727,  0.02786927,  0.08547455, -0.12870234,
                             -0.02555288,  0.11858013,  0.13305576,  0.07310204,  0.15149245,
                             -0.05708987, -0.064768 ,  0.18012332, -0.16880843, -0.27079877])
            alpha=0.0001
            N=len(X_train)
            grader_db(grad_x,grad_y,grad_w,grad_b)
```

Out[383]: True

```
In [384]: # prediction function used to compute predicted_y given the dataset X
def pred(w,b, X):
    N = len(X)
    predict = []
    for i in range(N):
        z=np.dot(w,X[i])+b
        predict.append(sigmoid(z))
    return np.array(predict)
```

Implementing logistic regression

```
In [385]: def train(X_train,y_train,X_test,y_test,epochs,alpha,eta0):
    ''' In this function, we will implement logistic regression'''
    train_loss = []
    test_loss = []
    w,b = initialize_weights(X_train[0]) # Initialize the weights

    #for every epoch
    for i in range(epochs):
        # for every data point(X_train,y_train)
        for j in range (N): #N=len(X_train)
            #computing gradient w.r.t w
            dw = gradient_dw(X_train[j],y_train[j],w,b,alpha,N)
            #computing gradient w.r.t b
            db = gradient_db(X_train[j],y_train[j],w,b)
            #updating w and b
            w = w+(eta0*dw)
            b = b+(eta0*db)

        #Predicting the output of X_train using pred function
        train_pred = pred(w,b, X_train)
        test_pred = pred(w,b, X_test)
        #computing the logloss for train and test data
        train_loss.append(logloss(y_train,train_pred))
        test_loss.append(logloss(y_test,test_pred))
    return w,b,train_loss,test_loss
```

```
In [386]: alpha=0.001
eta0=0.001
N=len(X_train)
epochs=20
w,b,train_loss,test_loss=train(X_train,y_train,X_test,y_test,epochs,alpha,eta0)
```

```
In [387]: #print thr value of weights w and bias b
print(w)
print(b)
```

```
[-0.41395269  0.19245258 -0.15005108  0.32635385 -0.22516783  0.58646754
 -0.42720461 -0.10027814  0.21483871  0.15555206  0.1788105  -0.01318643
 -0.06496816  0.36313959 -0.00985043]
-0.9016736323411028
```

```
In [388]: # these are the results we got after we implemented sgd and found the optimal weights
w=clf.coef_, b=clf.intercept_
```

```
Out[388]: (array([[ 0.00941423,  0.00697693, -0.00146073, -0.01509022, -0.01698113,
                   0.02630175,  0.02522022, -0.00619001,  0.00556551, -0.0252892 ,
                   -0.0182414 , -0.01740559,  0.01463554,  0.02461158, -0.03251764]]),
          array([-0.04853533]))
```

Goal of assignment

Compare your implementation and SGDClassifier's the weights and intercept, make sure they are as close as possible i.e difference should be in order of 10^{-2}

Grader function - 6

```
In [389]: #this grader function should return True
#the difference between custom weights and clf.coef_ should be less than or equal to 0.05
def difference_check_grader(w,b,coef,intercept):
    val_array=np.abs(np.array(w-coef))
    assert(np.all(val_array<=0.05))
    print('The custom weights are correct')
    return True
difference_check_grader(w,b,clf.coef_,clf.intercept_)
```

The custom weights are correct

```
Out[389]: True
```

Plot your train and test loss vs epochs

plot epoch number on X-axis and loss on Y-axis and make sure that the curve is converging


```
In [390]: import warnings
warnings.filterwarnings("ignore")

from matplotlib import pyplot as plt
epoch = list(range(1,(len(train_loss)+1),1))

plt.plot(epoch,train_loss)
plt.plot(epoch,test_loss)
plt.xlabel("epoch")
plt.ylabel("log loss")
plt.legend("Train loss","Test loss")
plt.title("train and test loss vs epochs")
plt.show
```

Out[390]: <function matplotlib.pyplot.show(*args, **kw)>

